

## 113-1 Operating System MP1

Group 45

111062332 朱誼學, 111062333 高英耀

Contributions	朱誼學	高英耀
Trace code	50%	50%
Implement	50%	50%
Report	50%	50%
Explenation	We've done all our work in discord vc	

## 1. Explain how system calls work in NachOS as requested in PartII-1

### a. SC\_Halt

SC\_Halt is called to shut down the whole process.

```
/* Stop Nachos, and print out performance stats */  
void Halt();
```

When Nachos starts up, it will call `void Machine::Run()`, and this function will never stop. Then, this function will call `void Machine::OneInstruction()`, when PCReg point to the address of Halt, SC\_Halt will be stored in r2 in start.S and than call `SYSCALL`:

```
.globl Halt  
.ent    Halt  
Halt:  
    addiu $2,$0,SC_Halt  
    syscall  
    j     $31  
.end    Halt
```

so `OneInstruction()` will call `RaiseException(SyscallException, 0);`

```
case OP_SYSCALL:  
    DEBUG(dbgTraCode, "In Machine::OneInstruction, RaiseException(SyscallException, 0), " << kernel->stats->totalTicks);  
    RaiseException(SyscallException, 0);  
    return;
```

Then `RaiseException` will call `ExceptionHandler(SyscallException)`, in `ExceptionHandler`, it will read r2

```
int type = kernel->machine->ReadRegister(2);
```

then execute SC\_Halt and call `SysHalt()`

```
case SC_Halt:  
    DEBUG(dbgSys, "Shutdown, initiated by user program.\n");  
    SysHalt();  
    cout << "in exception\n";  
    ASSERTNOTREACHED();  
    break;
```

`SysHalt()` then call `kernel->interrupt->Halt();`

```
void SysHalt() {  
    kernel->interrupt->Halt();  
}
```

Then delete the kernel to shut down nachos.

```
void Interrupt::Halt() {
#ifdef NO_HALT_STAT
    cout << "Machine halting!\n\n";
    cout << "This is halt\n";
    kernel->stats->Print();
#endif
    delete kernel; // Never returns.
}
```

## b. SC\_Create

Argument pointer in `Create()` will be stored in r4(MIPS conventions), `SC_Create` will be stored in r2 in `start.S` and then call `SYSCALL`, in `ExceptionHandler`, it will read r2 then execute `SC_Create`, also read r4 to get pointer and find the argument(filename) in memory with the pointer.

```
case SC_Create:
    val = kernel->machine->ReadRegister(4);
    {
        char *filename = &(kernel->machine->mainMemory[val]);
        // cout << filename << endl;
        status = SysCreate(filename);
        kernel->machine->WriteRegister(2, (int)status);
    }
    kernel->machine->WriteRegister(PrevPCReg, kernel->machine->ReadRegister(PCReg));
    kernel->machine->WriteRegister(PCReg, kernel->machine->ReadRegister(PCReg) + 4);
    kernel->machine->WriteRegister(NextPCReg, kernel->machine->ReadRegister(PCReg) + 4);
    return;
    ASSERTNOTREACHED();
    break;
```

Then passes the filename to `SysCreate`,

and `SysCreate()` passes to `FileSystem::Create()`, then `FileSystem::Create()` will Create a file in Disk.

```
int SysCreate(char *filename) {
    // return value
    // 1: success
    // 0: failed
    return kernel->fileSystem->Create(filename);
}
```

### c. SC\_PrintInt

Argument pointer in `PrintInt()` will be stored in `r4`(MIPS conventions), `SC_PrintInt` will be stored in `r2` in `start.S` and then call `SYSCALL`, in `ExceptionHandler`, it will read `r2` then execute `SC_PrintInt`, also read `r4` to get pointer and find the argument(int) in memory with the pointer.

```
case SC_PrintInt:
    DEBUG(dbgSys, "Print Int\n");
    val = kernel->machine->ReadRegister(4);
    DEBUG(dbgTraCode, "In ExceptionHandler(), into SysPrintInt, " << kernel->stats->totalTicks);
    SysPrintInt(val);
    DEBUG(dbgTraCode, "In ExceptionHandler(), return from SysPrintInt, " << kernel->stats->totalTicks);
    // Set Program Counter
    kernel->machine->WriteRegister(PrevPCReg, kernel->machine->ReadRegister(PCReg));
    kernel->machine->WriteRegister(PCReg, kernel->machine->ReadRegister(PCReg) + 4);
    kernel->machine->WriteRegister(NextPCReg, kernel->machine->ReadRegister(PCReg) + 4);
    return;
    ASSERTNOTREACHED();
    break;
```

Then passes the int value to `SysPrintInt`,

```
void SysPrintInt(int val) {
    DEBUG(dbgTraCode, "In ksyscall.h:SysPrintInt, into synchConsoleOut->PutInt, " << kernel->stats->totalTicks);
    kernel->synchConsoleOut->PutInt(val);
    DEBUG(dbgTraCode, "In ksyscall.h:SysPrintInt, return from synchConsoleOut->PutInt, " << kernel->stats->totalTicks);
}
```

and `SysPrintInt()` passes to `SynchConsoleOutput::PutInt()`, this function change the value to string, then passes the chars in the string to `SynchConsoleOutput::PutChar()` one by one.

```
void SynchConsoleOutput::PutInt(int value) {
    char str[15];
    int idx = 0;
    // sprintf(str, "%d\n", value); the true one
    sprintf(str, "%d\n", value); // simply for trace code
    lock->Acquire();
    do {
        DEBUG(dbgTraCode, "In SynchConsoleOutput::PutChar, into consoleOutput->PutChar, " << kernel->stats->totalTicks);
        consoleOutput->PutChar(str[idx]);
        DEBUG(dbgTraCode, "In SynchConsoleOutput::PutChar, return from consoleOutput->PutChar, " << kernel->stats->totalTicks);
        idx++;

        DEBUG(dbgTraCode, "In SynchConsoleOutput::PutChar, into waitFor->P(), " << kernel->stats->totalTicks);
        waitFor->P();
        DEBUG(dbgTraCode, "In SynchConsoleOutput::PutChar, return from waitFor->P(), " << kernel->stats->totalTicks);
    } while (str[idx] != '\0');
    lock->Release();
}
```

Then `SynchConsoleOutput::PutChar()` passes the char to `ConsoleOutput::PutChar(char ch)`.

```
void SynchConsoleOutput::PutChar(char ch) {
    lock->Acquire();
    consoleOutput->PutChar(ch);
    waitFor->P();
    lock->Release();
}
```

Then this function checks if hardware is ready for putchar, changes the state of the console to busy, schedules an interrupt and also displays the char.

```
void ConsoleOutput::PutChar(char ch) {
    ASSERT(putBusy == FALSE);
    WriteFile(writeFileNo, &ch, sizeof(char));
    putBusy = TRUE;
    kernel->interrupt->Schedule(this, ConsoleTime, ConsoleWriteInt);
}
```

**Interrupt::Schedule()** will record the **CallbackObject** and when to call back in the interrupt, and insert an interrupt to the pending list.

```
void Interrupt::Schedule(CallBackObj *toCall, int fromNow, IntType type) {
    int when = kernel->stats->totalTicks + fromNow;
    PendingInterrupt *toOccur = new PendingInterrupt(toCall, when, type);

    DEBUG(dbgInt, "Scheduling interrupt handler the " << intTypeNames[type] << " at time = " << when);
    ASSERT(fromNow > 0);

    pending->Insert(toOccur);
}
```

After the simulated time in kernel is skipped by the function **OneTick()** in **Run()**,

```
void Interrupt::OneTick() {
    MachineStatus oldStatus = status;
    Statistics *stats = kernel->stats;

    // advance simulated time
    if (status == SystemMode) {
        stats->totalTicks += SystemTick;
        stats->systemTicks += SystemTick;
    } else {
        stats->totalTicks += UserTick;
        stats->userTicks += UserTick;
    }
    DEBUG(dbgInt, "== Tick " << stats->totalTicks << " ==");
}
```

**OneTick()** will call **CheckIfDue()** to handle interrupts.

```
// check any pending interrupts are now ready to fire
ChangeLevel(IntOn, IntOff); // first, turn off interrupts
                             // (interrupt handlers run with
                             // interrupts disabled)
CheckIfDue(FALSE);         // check for pending interrupts
ChangeLevel(IntOff, IntOn); // re-enable interrupts
```

**CheckIfDue()** will call the **CallBack** function of the **CallbackObject**, which is **ConsoleOutput::CallBack()**.

```
inHandler = TRUE;
do {
    next = pending->RemoveFront(); // pull interrupt off list
    DEBUG(dbgTraCode, "In Interrupt::CheckIfDue, into callOnInterrupt->CallBack, " << stats->totalTicks);
    next->callOnInterrupt->CallBack(); // call the interrupt handler
    DEBUG(dbgTraCode, "In Interrupt::CheckIfDue, return from callOnInterrupt->CallBack, " << stats->totalTicks);
    delete next;
} while (!pending->IsEmpty() && (pending->Front()->when <= stats->totalTicks));
inHandler = FALSE;
return TRUE;
```

`ConsoleOutput::CallBack()` will update the state of hardware and call `SynchConsoleOutput::CallBack()`

```
void ConsoleOutput::CallBack() {  
    DEBUG(dbgTraCode, "In ConsoleOutput::CallBack(), " << kernel->stats->totalTicks);  
    putBusy = FALSE;  
    kernel->stats->numConsoleCharsWritten++;  
    callWhenDone->CallBack();  
}
```

`SynchConsoleOutput::CallBack()` is to tell nachos it is ready to print next char.

```
void SynchConsoleOutput::CallBack() {  
    DEBUG(dbgTraCode, "In SynchConsoleOutput::CallBack(), " << kernel->stats->totalTicks);  
    waitFor->V();  
}
```

## 2. Explain our implementation as requested in PartII-2

- Working items:
  - (a). `OpenFileId Open(char *name);`  
Open a file with the name, and return its corresponding `OpenFileId`.  
**Return -1 if it fails to open the file.**
  - (b). `int Write(char *buffer, int size, OpenFileId id);`  
Write “size” characters from the buffer into the file, and return the number of characters actually written to the file.  
**Return -1, if it fails to write the file.**
  - (c). `int Read(char *buffer, int size, OpenFileId id);`  
Read “size” characters from the file to the buffer, and return the number of characters actually read from the file.  
**Return -1, if it fails to read the file.**
  - (d). `int Close(OpenFileId id);`  
Close the file with id.  
**Return 1 if successfully close the file. Otherwise, return -1.**  
**Need to delete the OpenFile after you close the file**

### Explanation:

In **test/start.s**, we add a few assembly code for `Open`, `Write`, `Read`, `Close`, so that the function define in `syscall.h` can link to `start.S`, and then be translated into assembly code.

```
.globl Open
.ent    Open
Open:
    addiu $2, $0, SC_Open
    syscall
    j     $31
.end    Open
```

In **userprog/syscall.h**, we uncommented the constant define, which is also the offset for 'addiu' in `start.S`

```
#define SC_Open 6
#define SC_Read 7
#define SC_Write 8
#define SC_Seek 9
#define SC_Close 10
```

In **userprog/exception.cc**, we add 4 new cases into the switch structure.

**SC\_Open:** Read the address of filename from register 4 and use it to get the filename from mainMemory, then call `OpenFileID SysOpen(char *name)`, after that save the returned file id into register2, and update the PC register at the end.

```
case SC_Open:
    val = kernel->machine->ReadRegister(4);
    {
        char *filename = &(kernel->machine->mainMemory[val]);
        fileID = SysOpen(filename);
        kernel->machine->WriteRegister(2, (int)fileID);
    }
    kernel->machine->WriteRegister(PrevPCReg, kernel->machine->ReadRegister(PCReg));
    kernel->machine->WriteRegister(PCReg, kernel->machine->ReadRegister(PCReg) + 4);
    kernel->machine->WriteRegister(NextPCReg, kernel->machine->ReadRegister(PCReg) + 4);
    return;
    ASSERTNOTREACHED();
    break;
```

**SC\_Read:** Get the attributes from register 4, 5, 6 respectively. gain the buffer content from mainMemory by address, and pass all the attributes into `int SysRead(char *buffer, int size, OpenFileId id)`. The returned status shows how many characters are read into the buffer. Likely, then save the value into register 2, and update the PC register.

```
case SC_Read:
    val = kernel->machine->ReadRegister(4);
    numChar = kernel->machine->ReadRegister(5);
    fileID = kernel->machine->ReadRegister(6);
    {
        char *buf = &(kernel->machine->mainMemory[val]);
        status = SysRead(buf, numChar, fileID);
        kernel->machine->WriteRegister(2, (int)status);
    }
    kernel->machine->WriteRegister(PrevPCReg, kernel->machine->ReadRegister(PCReg));
    kernel->machine->WriteRegister(PCReg, kernel->machine->ReadRegister(PCReg) + 4);
    kernel->machine->WriteRegister(NextPCReg, kernel->machine->ReadRegister(PCReg) + 4);
    return;
    ASSERTNOTREACHED();
    break;
```



**SC\_Write:** Same as SC\_Read but using `int SysWrite(char *buffer, int size, OpenFileId id)`. And the status stands for how many characters are written into the buffer instead of read.

```
case SC_Write:
    val = kernel->machine->ReadRegister(4);
    numChar = kernel->machine->ReadRegister(5);
    fileID = kernel->machine->ReadRegister(6);
    {
        char *buf = &(kernel->machine->mainMemory[val]);
        status = SysWrite(buf, numChar, fileID);
        kernel->machine->WriteRegister(2, (int)status);
    }
    kernel->machine->WriteRegister(PrevPCReg, kernel->machine->ReadRegister(PCReg));
    kernel->machine->WriteRegister(PCReg, kernel->machine->ReadRegister(PCReg) + 4);
    kernel->machine->WriteRegister(NextPCReg, kernel->machine->ReadRegister(PCReg) + 4);
    return;
    ASSERTNOTREACHED();
    break;
```

**SC\_Close:** Read the file id from register 4 and save into val, passing it into `int SysClose(OpenFileId id)`, and get the returned value called success. Later, save the result into register 2. At last, update the PC register as well.

```
case SC_Close:
    val = kernel->machine->ReadRegister(4);
    {
        int success;
        success = SysClose(val);
        kernel->machine->WriteRegister(2, (int)success);
    }
    kernel->machine->WriteRegister(PrevPCReg, kernel->machine->ReadRegister(PCReg));
    kernel->machine->WriteRegister(PCReg, kernel->machine->ReadRegister(PCReg) + 4);
    kernel->machine->WriteRegister(NextPCReg, kernel->machine->ReadRegister(PCReg) + 4);
    return;
    ASSERTNOTREACHED();
    break;
```

In **userprog/ksyscall.h**:

Define the functions that we called in **exception.cc**, which simply pass the return value from **kernel->fileSystem** to **exceptionHandler**.

```
OpenFileId SysOpen(char *name)
{
    return kernel->fileSystem->OpenAFile(name);
}

int SysRead(char *buffer, int size, OpenFileId id){
    return kernel->fileSystem->ReadFile(buffer, size, id);
}

int SysWrite(char *buffer, int size, OpenFileId id){
    return kernel->fileSystem->WriteFile(buffer, size, id);
}

int SysClose(OpenFileId id){
    return kernel->fileSystem->CloseFile(id);
}
```

In **userprog/filesys.h**:

We have to define a 2d char array in order to store the opened file's name.

```
// An array to memorize the name of opened file.
char OpenFileName[20][100];
```

In **OpenAFile**, at first we will call a predefined function **inline OpenFile \*FileSystem::Open(char \*name)** which will return an **OpenFile** instance. If it fails to open a file, **OpenAFile** will return -1. Moreover, we adopt a brute force search for handling duplicate file opening. If it matches the same filename, return -1. Next, we parse over the **OpenFileTable**, finding an index with nothing is in the corresponding memory space. When found, put the **OpenFile** instance in, and memorize the filename. Then return the index.

```
OpenFileId OpenAFile(char *name) {
    OpenFile* f = Open(name);
    if(f == NULL) return -1;
    for(int i=0; i<20; i++)
        if(strcmp(OpenFileName[i], name) == 0) return -1;
    int ans = -1;
    for(int i=0; i<20; i++)
        if(OpenFileTable[i] == NULL){
            OpenFileTable[i] = f;
            strcpy(OpenFileName[i], name);
            ans = i;
            break;
        }
    return (OpenFileId)ans;
}
```

In `WriteFile`, we get the corresponding file with file id as the index. If it fails to open the file or id is out of range, the function returns -1, else calls the predefined function `inline int OpenFile::Write(char *from, int numBytes)`, this function returns the number of written characters.

```
int WriteFile(char *buffer, int size, OpenFileId id){
    OpenFile* f = OpenFileTable[id];
    if(f != NULL && id >= 0 && id < 20){
        return f->Write(buffer, size);
    }
    else return -1;
}
```

`ReadFile` is similar to `WriteFile` but called `inline int OpenFile::Read(char *into, int numBytes)` instead of `f->Write()`.

```
int ReadFile(char *buffer, int size, OpenFileId id){
    OpenFile* f = OpenFileTable[id];
    if(f != NULL && id >= 0 && id < 20){
        return f->Read(buffer, size);
    }
    else return -1;
}
```

In `CloseFile`, we get the file we're going to close by its file id. Also, check whether we get the `OpenFile` pointer successfully, and the id is also in `[0, 20)`. Then, clear the corresponding row in the name array. Delete the file by its pointer and return 1.

```
int CloseFile(OpenFileId id){
    OpenFile* f = OpenFileTable[id];
    if(f != NULL && id >= 0 && id < 20){
        OpenFileTable[id] = NULL;
        for(int i=0; i<100; i++) OpenFileName[id][i] = '\0';
        delete f;
        return 1;
    }
    else return -1;
}
```

### 3. The difficulties we encounter when implementing this assignment:

高英耀: At the very beginning, I did not know where to start, and I was afraid that changing something would cause the crash of nachos. So I asked Chatgpt for an example, and got an answer which seems to be very reliable. Anyway, now I have implemented some functions, so I should be less afraid to do some implementations in future, even though no one told me how to do it.

朱誼學: When I am tracing the codes, I am very confused about why, there are many function such as halt() didn't be defined but still callable. After reading the problems and responses in the discussion list in eeclass, I have totally understood how start.S and the files in userprog interact with each other. Moreover, there isn't any TODO labeled for us, so I took a little bit more time to figured out where should I start my implementation.

### 4. Feedback:

高英耀: 老師教得很好, 上課的內容都能使我輕易的理解與吸收, 認真上課使我在trace code時能順利理解, 也感謝討論區同學與李秉綸助教在討論區的提問語回答, 在實作的過程遇到的問題大多都能在討論區找到答案, 李秉綸助教的回答都十分熱情與迅速, 值得五星好評, 實在感謝這堂OS的教授與助教們, 為我們清華CS的學生建造出一個十分良好的學習環境。

朱誼學: Thanks for my classmates and TAs asking and answering question in eeclass, so that I can find the solution of most of my problems in there. And I really learn a lot in the professor's lesson and in this lab. Combining the detailed teaching and practical implementation, I feel like the Operating system is not that abstract anymore.